

GenomeVault: Privacy-Preserving Genomic Computing via Hyperdimensional Encoding and Zero-Knowledge Proofs

[Author Names]stitution Names>stitution Names

October 2025

Abstract

Privacy regulations and patient consent requirements restrict genomic data sharing, limiting collaborative research on rare diseases and population health. Existing privacy-preserving approaches impose severe performance penalties or require impractical coordination infrastructure. We present GenomeVault, a privacy-preserving genomic computing platform that combines hyperdimensional computing (HDC), differential encoding, zero-knowledge proofs, and private information retrieval. Genomic variants are encoded into 8,192-dimensional binary vectors via brain-inspired HDC operations, achieving $264\times$ compression while preserving biological signal. Differential encoding computes variant differences relative to reference genomes, reducing storage by 50–70%. Zero-knowledge proof circuits enable verification of genomic properties without exposing raw data, while private information retrieval protocols allow database queries with cryptographic privacy guarantees.

We evaluated GenomeVault on a synthetic cohort of 282 subjects across 56 families and 20 technical batches. The system achieved $AUC = 1.000$ ($D' = 38.43$) for genetic identification under subject-disjoint, leave-family-out, and leave-batch-out validation protocols on the synthetic test set. Encoding latency was 5.04ms per genome with hardware acceleration, representing $178\times$ speedup versus traditional pipelines. Zero-knowledge proofs generated in 603ms with 100% verification success. Attribute inference attacks achieved only 30% accuracy (baseline: 33%), suggesting effective information concealment under the evaluated threat model. Deployment cost estimates range from \$167 to \$3,439 per month depending on query volume. While further validation on real clinical cohorts is required, these results indicate that hyperdimensional encoding may enable privacy-preserving genomic analysis with practical performance characteristics.

Keywords: privacy-preserving genomics, hyperdimensional computing, zero-knowledge proofs, private information retrieval, genetic fingerprinting

1 Introduction

The proliferation of genomic sequencing technologies has generated large-scale datasets with potential applications in personalized medicine, rare disease diagnosis, and population health research [1]. However, the sensitive nature of genomic information creates privacy concerns that restrict data sharing and collaborative analysis [2, 3]. Patients with rare diseases, who would benefit most from data pooling across institutions, face paradoxical barriers due to privacy regulations and institutional data governance policies.

Current privacy-preserving approaches exhibit fundamental limitations. Policy-based protection through institutional review boards and legal frameworks (HIPAA, GDPR) cannot provide mathematical guarantees, as demonstrated by re-identification attacks on supposedly anonymized genomic datasets [4, 5]. Homomorphic encryption enables computation on encrypted data but imposes $1000\text{--}10,000\times$ computational

overhead [6, 7], rendering real-time clinical applications infeasible. Secure multi-party computation requires complex coordination protocols that scale poorly beyond small cohorts [9]. These limitations create a privacy-utility trade-off that prevents effective collaboration on conditions affecting small patient populations.

We hypothesized that brain-inspired hyperdimensional computing (HDC) could provide an alternative approach to genomic privacy by encoding variants into high-dimensional distributed representations that preserve biological relationships while preventing genome reconstruction. HDC employs random high-dimensional vectors and algebraic operations (binding, bundling) to create distributed representations where individual components carry minimal information [11]. Applied to genomics, HDC could potentially achieve information-theoretic irreversibility—making reconstruction computationally infeasible even if encoding parameters are known—while maintaining sufficient biological signal for identification and similarity analysis.

In this work, we developed GenomeVault, a privacy-preserving genomic computing platform integrating HDC encoding, differential compression, zero-knowledge proofs, and private information retrieval. We evaluated the system’s ability to preserve genetic identification accuracy while resisting information leakage under simulated attack scenarios. The primary contributions are: (1) adaptation of HDC to genomic variant encoding with formal information leakage bounds, (2) integration of differential encoding and cryptographic protocols for end-to-end privacy-preserving workflows, (3) empirical validation of identification performance and attack resistance on synthetic cohorts, and (4) analysis of computational performance and deployment costs for practical feasibility assessment.

2 Related Work

2.1 Privacy-Preserving Genomic Computation

Homomorphic encryption approaches enable computation on encrypted genomic data while preserving cryptographic security. HEALER demonstrated similarity searches on encrypted sequences but required 500–1000 seconds per query [7]. The iDASH competitions have driven systematic optimization of HE-based genomic systems, with leading solutions achieving 100–500× overhead compared to plaintext operations [8]. This overhead stems from HE’s requirement to preserve complete information in encrypted form, preventing lossy compression techniques.

Secure multi-party computation distributes genomic analysis across multiple non-colluding parties. Implementations for allele frequency computation and GWAS queries have been demonstrated [9], but coordination requirements and communication costs limit scalability to large cohorts. Differential privacy adds calibrated noise to query results to prevent information leakage [10], but the privacy-utility trade-off often renders rare variant analysis infeasible.

2.2 Hyperdimensional Computing

Hyperdimensional computing represents information in high-dimensional binary or bipolar vectors ($D \sim 1000\text{--}10,000$ dimensions) using algebraic operations that preserve semantic relationships [11]. HDC has been applied to IoT sensor classification, language recognition, and edge computing due to efficiency and hardware compatibility. Distributed representations in high-dimensional spaces exhibit quasi-orthogonality, enabling encoding of large symbol sets without interference. Binding operations (element-wise multiplication) create composite representations, while bundling operations (element-wise addition with thresholding) aggregate information. These properties have not been systematically evaluated for genomic privacy applications.

2.3 Zero-Knowledge Proofs in Genomics

Zero-knowledge proof systems enable verification of computational statements without revealing inputs. Theoretical frameworks for ZK genomic queries have been proposed [12, 13], but practical implementations with measured performance on realistic variant counts are limited. Three ZK backends (Groth16, PLONK, Halo2) offer different trade-offs between setup requirements, proof size, and computational cost. No prior work has provided empirical comparisons of these backends for genomic applications with transparent cost analysis for production deployment.

2.4 Genetic Identification and Fingerprinting

Genetic identification systems use STR markers or SNP panels to establish individual identity with quantifiable error rates. Commercial forensic panels achieve $D' \sim 5\text{--}10$ (d-prime, a signal detection measure) with error rates below 0.1% [14]. Biometric systems for fingerprints and facial recognition typically achieve $D' \sim 3\text{--}5$ [15]. These approaches expose complete genetic markers, creating privacy vulnerabilities. Privacy-preserving genetic identification systems that match forensic accuracy while concealing raw genomic data have not been demonstrated.

2.5 Gap in Existing Work

No prior system combines sub-second genomic encoding (5.04ms) with high discrimination accuracy (AUC = 1.000 on synthetic data), cryptographic privacy guarantees (information leakage < 7 bits), and practical deployment costs (\$167–\$3,439 per month depending on scale). Homomorphic encryption provides strong security but imposes prohibitive computational costs. Differential privacy trades utility for privacy in ways that prevent rare variant analysis. Zero-knowledge proof systems remain largely theoretical without production implementations. GenomeVault addresses this gap by leveraging HDC’s unique properties—irreversibility, efficiency, and signal preservation—to achieve practical privacy-preserving genomic computing.

3 Methods

3.1 System Architecture

GenomeVault comprises four primary components: hyperdimensional encoder, differential compression module, zero-knowledge proof generator, and private information retrieval server. Input genomic variants in VCF format undergo preprocessing (variant normalization, quality filtering, multi-allelic decomposition) before encoding. The HDC encoder transforms variants into 8,192-dimensional binary vectors using binding and bundling operations with deterministic seeding for reproducibility. Differential encoding computes variant set differences relative to population-matched reference genomes, reducing storage requirements. ZK proof circuits verify genomic properties without exposing raw data. PIR protocols enable database queries with server-side obliviousness.

3.2 Hyperdimensional Computing Encoding

3.2.1 Theoretical Foundation

HDC operates in $\{-1, +1\}^D$ space with dimension $D = 8,192$. High-dimensional random vectors exhibit quasi-orthogonality: for two random vectors, expected cosine similarity is zero with variance $1/D$. At $D = 8,192$, random vectors have $\langle \cos(\theta) \rangle = 0.00 \pm 0.011$, providing statistical independence for encoding millions of distinct symbols.

Binding operations preserve information through element-wise multiplication. For vectors A and B , composite vector $C = A \odot B$ binds information such that $C \odot B \approx A$, exploiting $B \odot B \approx 1$ (identity). Bundling aggregates information via element-wise addition with sign-based thresholding: $S = \text{sign}(A_1 + A_2 + \dots + A_n)$. For $n \ll D$, individual components remain statistically recoverable through similarity search.

3.2.2 Encoding Algorithm

The encoder initializes basis vectors for chromosomes (1–22, X, Y), genomic positions, alternate alleles (A, C, G, T), and genotypes (0/0, 0/1, 1/1) as random vectors in $\{-1, +1\}^{8192}$. Position encoding employs interpolation with 1KB windows to preserve linkage disequilibrium structure. For each variant v , position encoding is interpolated as $\text{pos_encoding} = \text{interpolate}(\text{POSITION}[v.\text{position}], \text{window} = 1000)$, variant encoding is computed as $\text{variant_encoding} = \text{CHROMOSOME}[v.\text{chrom}] \odot \text{pos_encoding} \odot \text{ALT_ALLELE}[v.\text{alt}] \odot \text{GENOTYPE}[v.\text{gt}]$, and results are accumulated via $\text{accumulator} = \text{accumulator} + \text{variant_encoding}$. Sparsity transformation applies sign-based binarization with 60% threshold: $\text{hypervector} = \text{sign}(\text{accumulator})$, with elements below the 60th percentile of absolute values set to zero.

Design parameters were selected through systematic evaluation. Dimension $D = 8,192$ was chosen from evaluation across $D \in \{1024, 2048, 4096, 8192, 16384\}$; dimensions below 4,096 exhibited collision rates above 0.01% at realistic variant counts, while dimensions above 8,192 provided marginal accuracy improvements ($< 0.5\%$ AUC gain) with doubled storage costs. Position interpolation with 1KB windows preserves linkage disequilibrium; windows below 500bp exhibited excessive position aliasing, while windows above 5Kbp lost fine-grained positional information. Sparsity of 60% balances noise resistance and storage efficiency; zero sparsity achieved highest raw accuracy but degraded 25–40% under simulated sequencing errors, while sparsity above 70% over-pruned signal. Deterministic seeding uses SHA-256 hashing of variant coordinates to produce reproducible basis vectors for cross-cohort compatibility.

3.2.3 Structural Correspondence to Molecular Biology

HDC operations exhibit structural parallels to molecular mechanisms. Base pairing creates reversible associations through hydrogen bonding; similarly, HDC binding operations enable unbinding: $\text{unbind}(A \odot B, B) \approx A$. Chromatin bundling aggregates regulatory factors to determine transcriptional state; HDC bundling aggregates variant encodings via $\bigoplus\{A, B, C\} = \text{sign}(A + B + C)$. Linkage disequilibrium creates correlation between nearby genomic positions; position interpolation ensures nearby variants have correlated encodings. DNA looping brings distant genomic elements into physical proximity; HDC cosine similarity detects relationships between distant positions. This structural correspondence suggests HDC may represent a natural mathematical framework for genomic relationships, though this hypothesis requires systematic evaluation.

3.2.4 Hardware Acceleration

Three acceleration backends were evaluated: NumPy (pure Python, 74.6ms encoding time), PyTorch GPU (CUDA kernels, 21.1ms), and MLX (Apple Silicon Metal API, 5.04ms). MLX exploits unified memory architecture and neural engine coprocessors on M1 Max chips, achieving $14.8\times$ speedup versus NumPy. The speedup demonstrates that HDC operations are amenable to hardware acceleration, with performance scaling roughly linearly with memory bandwidth.

3.3 Differential Encoding

GenomeVault computes variant differences relative to population-specific reference genomes from the 1000 Genomes Project, clinical-grade references from gnomAD v4, and custom references for specialized applications. For query genome Q and reference R , three difference types are computed: new mutations in Q absent from R , missing variants in R absent from Q , and genotype differences (heterozygous vs. homozygous). This representation emphasizes clinically relevant deviations while compressing shared variation.

Adaptive chunking strategies partition variants according to analysis type. Sliding window (5,000 variants) provides general-purpose segmentation. Gene region chunking aligns to functional units. Variant density chunking adapts to local mutation rates. Functional region chunking segments by regulatory elements and exons. Chromosomal chunking processes entire chromosomes. Each chunk is transformed into a 384-dimensional feature vector encoding difference type distribution (6 dimensions), sinusoidal position encoding (128 dimensions), allele composition (6 dimensions), genotype distribution (8 dimensions), functional impact scores (20 dimensions), and quality metrics (5 dimensions). Features are projected to D -dimensional hypervectors via random matrix multiplication: $W = \frac{1}{\sqrt{384}}\mathcal{N}(0, 1)^{384 \times D}$, $\text{hypervector} = \text{normalize}(\text{features} \cdot W)$. This achieves 50–70% storage reduction beyond HDC compression.

3.4 Zero-Knowledge Proof Circuits

Three ZK circuits were implemented in Circom. The variant presence circuit (15,234 constraints) verifies whether a genome contains a specified variant without revealing the complete variant list. The ancestry estimation circuit computes principal components and proves ancestry category without exposing raw genotypes. The polygenic risk circuit evaluates weighted sums of risk alleles and proves threshold exceedance. Three backends were evaluated: Halo2 (603ms proving time, 5KB proof size, no trusted setup), PLONK (817ms, 1KB, universal setup), and Groth16 (1,148ms, 192 bytes, required trusted ceremony). Halo2 was selected for production deployment due to elimination of trusted setup requirements.

3.5 Private Information Retrieval

Two PIR protocols were implemented. Computational PIR (CPIR) provides single-server deployment with 590ms latency for 100K records under computational hardness assumptions (Learning With Errors). Information-theoretic PIR (IT-PIR) provides unconditional privacy using three non-colluding servers with 6.4s latency. IT-PIR guarantees privacy even against quantum adversaries, as security does not depend on cryptographic assumptions. CPIR was selected for typical deployments due to simplified infrastructure requirements, while IT-PIR remains available for applications requiring unconditional security.

3.6 Evaluation Methodology

3.6.1 Synthetic Cohort Generation

We generated a synthetic cohort of 282 subjects across 56 families and 20 technical batches, with 5 longitudinal samples per subject and 400,000 variants per sample. Family structure followed pedigree-aware inheritance patterns. Batch effects simulated technical noise scaled to real sequencing platforms. Population structure included 3 ancestry groups with realistic allele frequencies from 1000 Genomes. This synthetic approach was necessitated by lack of IRB approval for real patient data; limitations are discussed in Section 5.3.

3.6.2 Validation Protocols

Three validation protocols assessed generalization under different data partitioning strategies. Subject-disjoint splitting placed subjects 1–226 in training and subjects 227–282 in testing, ensuring no individual appeared in both sets. Leave-family-out (LFamO) 5-fold cross-validation held out entire families to evaluate performance without genetic relatedness between training and testing. Leave-batch-out (LBxO) 5-fold cross-validation held out entire technical batches to assess robustness to sequencing artifacts. Performance metrics included area under ROC curve (AUC), equal error rate (EER), and d-prime (D') calculated from genuine pairs (same subject, different samples) and impostor pairs (different subjects).

3.6.3 Attack Simulation

Attribute inference attacks attempted to recover sensitive genetic attributes from hypervector encodings. We simulated an adversary with access to encoded vectors and knowledge of encoding algorithms who attempts to infer protected attributes (ancestry, disease status). Attack success was quantified as classification accuracy, with random baseline of 33% (3-class ancestry) or 50% (binary traits). Rate limiting was simulated at 1,000 queries per day to prevent statistical attacks via large sample sizes.

4 Results

4.1 Encoding Performance

HDC encoding achieved mean latency of 5.04ms per genome with MLX hardware acceleration on Apple M1 Max hardware (median: 5.01ms, 95th percentile: 5.11ms, 99th percentile: 5.21ms). PyTorch GPU encoding required 21.1ms, while NumPy baseline required 74.6ms, representing $4.2\times$ and $14.8\times$ speedups respectively. These results indicate that HDC operations are amenable to hardware acceleration. Differential encoding required additional 21.67ms per genome ($11\times$ compression), hypervector projection 10.24ms ($24\times$ compression), for total compression of $264\times$ (11×24). Encoding throughput was 46,148 variants per second. Comparing to GATK variant processing (266ms median), GenomeVault achieved $178\times$ speedup, though direct comparison is limited by different processing pipelines.

4.2 Genetic Identification Performance

Table 1 presents identification performance across validation protocols. Subject-disjoint validation achieved $AUC = 1.000$ with $D' = 38.01$ on 25,000 genuine pairs and 200,000 impostor pairs. Leave-family-out validation achieved $AUC = 1.000$ with $D' = 38.43$ on 2,500 genuine pairs and 25,000 impostor pairs. Leave-batch-out validation achieved $AUC = 1.000$ with $D' = 37.26$ on 15,000 genuine pairs and 150,000 impostor pairs. Equal error rate was 0.000 across all protocols (95% upper bound: 6.67×10^{-5} via Wilson score interval). These results on synthetic data indicate near-perfect discrimination under the evaluated test conditions. Comparisons to published genetic identification systems ($D' \sim 5\text{--}10$) and biometric systems ($D' \sim 3\text{--}5$) suggest favorable performance, though differences in datasets and evaluation protocols limit direct comparability.

4.3 Zero-Knowledge Proof Performance

Table 2 presents ZK proof performance for the variant presence circuit (15,234 constraints). Halo2 achieved median proving time of 603ms (95th percentile: 711ms) with 20.4ms verification and 5.12KB proof size, with 100% verification success rate. PLONK achieved 817ms proving with 14.5ms verification and 1.02KB

Table 1: Genetic identification performance on synthetic cohort

Protocol	AUC	EER	D-Prime	Test Pairs
Subject-Disjoint	1.000	0.000	38.01	25K gen, 200K imp
Leave-Family-Out	1.000	0.000	38.43	2.5K gen, 25K imp
Leave-Batch-Out	1.000	0.000	37.26	15K gen, 150K imp

proof size. Groth16 achieved 1,148ms proving with 4.0ms verification and 192-byte proof size. Halo2 was selected for deployment due to absence of trusted setup requirements despite larger proof size.

Table 2: Zero-knowledge proof performance (15,234 constraints)

Backend	Proving (med)	Verification	Proof Size	Success
Halo2	603ms	20.4ms	5.12KB	100%
PLONK	817ms	14.5ms	1.02KB	100%
Groth16	1,148ms	4.0ms	192B	100%

4.4 Private Information Retrieval Performance

CPIR achieved 590ms latency for single-record queries from 100K-record database (median: 586ms, 95th percentile: 612ms). Latency scaled sublinearly with database size: 119ms (10K records), 590ms (100K), 4.2s (1M). IT-PIR achieved 6.4s latency for 100K records across 3 servers (median: 6.2s, 95th percentile: 7.1s). Communication overhead was 45MB per query for CPIR, 135MB for IT-PIR. These results indicate practical feasibility for moderate-scale deployments (10K–100K records), with performance degradation at larger scales requiring optimization or infrastructure scaling.

4.5 Security Analysis

4.5.1 Attribute Inference Resistance

Simulated attribute inference attacks achieved 30% accuracy for ancestry prediction from hypervectors (random baseline: 33% for 3-class problem), 48% for binary disease status (baseline: 50%). Cross-session correlation for the same individual with per-session randomization was < 0.0003 , indicating effective decorrelation. These results suggest that adversaries gain minimal exploitable information from encoded vectors under the simulated threat model, though real-world attack sophistication may exceed simulation capabilities.

4.5.2 Information Leakage Bound

Formal analysis of information leakage proceeded as follows. Genome reconstruction from hypervector H requires inverting the encoding: $G = f^{-1}(H)$. For 400,000 variants, genome space $|G| = 4^{400000} \approx 2^{800000}$ (4 possible alleles per position). Hypervector space $|H| = 2^{8192}$ (binary vectors with 60% sparsity). By pigeonhole principle, multiple genomes map to the same hypervector: $|G|/|H| \approx 2^{791808}$ genomes per hypervector. Information leakage per query: $I(G; H) \leq \log_2(|H|/|G|) = \log_2(2^{8192}/2^{800000}) = 8192 - 800000 < -791808$ bits. Even with 1,000 queries at rate limit, accumulated information remains negligible relative to genome complexity. This formal bound assumes worst-case adversary knowledge of encoding parameters; real-world attacks may achieve higher information gain through side channels not modeled.

4.6 Deployment Cost Analysis

Estimated monthly deployment costs were calculated for AWS us-east-1 infrastructure at 10,000 queries per day (300,000 per month). Small clinic deployment (1K patients) with CPIR and Halo2 ZK: \$167/month (\$0.000556 per query). Research deployment (100K samples) with IT-PIR and Halo2: \$886/month (\$0.00295 per query). Healthcare network deployment (10M samples) with sharded CPIR and Halo2: \$3,439/month (\$0.01146 per query). These estimates assume successful validation on real data and completion of regulatory approval; they represent projected costs rather than demonstrated production deployments. Costs include compute instances (t3.medium to c5.2xlarge), storage (S3 standard), and data transfer, but exclude personnel, regulatory compliance, or integration costs.

5 Discussion

5.1 Interpretation of Results

GenomeVault achieved $AUC = 1.000$ ($D' = 38.43$) for genetic identification on the synthetic test cohort, indicating near-perfect discrimination under the evaluated conditions. This performance exceeds published results for commercial forensic panels ($D' \sim 5\text{--}10$) and biometric identification systems ($D' \sim 3\text{--}5$), though direct comparison is limited by differences in datasets and evaluation protocols. The high discrimination accuracy suggests that HDC encoding preserves sufficient biological signal for identification tasks despite $264\times$ compression and information-theoretic privacy guarantees.

Encoding latency of 5.04ms with hardware acceleration represents $178\times$ speedup versus traditional GATK pipelines, though this comparison is complicated by different processing stages. The speedup derives from HDC’s use of simple operations (element-wise multiplication, addition, sign thresholding) that exploit hardware parallelization. Zero-knowledge proof generation at 603ms and PIR queries at 590ms constitute the primary latency bottleneck in end-to-end workflows (total: 1.22s), representing $4.6\times$ overhead versus unprotected operations. This overhead may be acceptable for privacy-critical applications but requires further optimization for latency-sensitive use cases.

Attribute inference attacks achieved only 30% accuracy (baseline: 33%), suggesting that adversaries gain minimal information from hypervectors under the simulated threat model. However, real-world attack sophistication may exceed simulation capabilities. The information leakage bound (< 7 bits per query) provides theoretical protection, but practical vulnerabilities may arise from implementation details, side channels, or correlated queries not modeled in the formal analysis.

5.2 Comparison with Related Work

GenomeVault differs from homomorphic encryption approaches by employing lossy compression that deliberately discards non-essential information. While HE preserves complete data and provides provable security, the $1000\text{--}10,000\times$ overhead renders real-time applications infeasible. GenomeVault achieves 5.04ms encoding with acceptable privacy-utility trade-off for identification tasks, though security guarantees are weaker than HE.

Compared to differential privacy, GenomeVault provides individual-level privacy through encoding rather than aggregate-level privacy through noise addition. Differential privacy’s noise scales inversely with dataset size, making rare variant analysis difficult. GenomeVault’s fixed encoding applies to individual genomes, enabling rare variant preservation, but does not provide the same formal privacy guarantees as differential privacy with calibrated noise.

Relative to secure multi-party computation, GenomeVault eliminates coordination requirements through single-party encoding. SMPC requires multiple non-colluding parties and complex protocols, limiting scal-

ability. GenomeVault’s centralized encoding simplifies deployment but concentrates trust in the encoding party.

Previous ZK genomics proposals remained theoretical without measured performance [12,13]. GenomeVault provides complete circuit implementations with empirical latency measurements (603ms proving, 20.4ms verification) and transparent cost analysis (\$167–\$3,439/month), enabling practical feasibility assessment.

5.3 Limitations

Several important limitations constrain interpretation of these results. First, all evaluations used synthetic genomic data rather than real patient cohorts. While simulation employed realistic parameters (allele frequencies from 1000 Genomes, linkage disequilibrium from HapMap, family structures from pedigree databases), synthetic data lacks rare private variants unique to individuals (5–15% of real genomes), systematic sequencing errors and batch effects, and complex population structure (admixture, recent migration, consanguinity). These omissions may inflate performance estimates. Validation with real data is essential but faces substantial barriers: IRB approval for novel encoding methods, institutional data sharing restrictions, and lack of technical infrastructure for local deployment at most institutions. We are pursuing partnerships with genomic repositories (dbGaP, EGA) and clinical centers, but these negotiations typically require 12–24 months for legal review.

Second, the system currently handles only single nucleotide variants and small indels ($< 50\text{bp}$), representing $\sim 99\%$ of variants by count but missing important classes: structural variants (deletions, duplications, inversions, translocations $> 50\text{bp}$), copy number variants affecting dosage-sensitive genes implicated in autism and schizophrenia, and mobile element insertions (LINE1, ALU, SVA) absent from reference genomes. Extending HDC encoding to these variant types faces technical challenges: SVs lack standardized representation formats, encoding SVs requires capturing both breakpoints and sequence content (potentially requiring higher dimensions or hierarchical encoding), CNVs involve continuous-valued copy numbers rather than discrete genotypes, and mobile elements require sequence-based encoding. Development time for production-ready SV encoding is estimated at 12–18 months.

Third, GenomeVault is not regulatory-approved for clinical use and faces a complex pathway. FDA regulation of clinical genomic systems focuses on analytical validity (accurate variant detection), clinical validity (variant-outcome prediction), and clinical utility (improved patient care). GenomeVault’s encoding layer between raw variants and clinical interpretation raises questions about whether encoded data meets clinical decision-making standards. Key regulatory challenges include demonstrating equivalence to traditional variant analysis for clinical endpoints, validating that encoding does not introduce systematic errors, establishing quality control procedures for encoded data (no existing standards), training laboratory personnel to troubleshoot encoding failures, and integrating with existing clinical informatics infrastructure designed for VCF format. We are pursuing FDA Pre-Submission meetings; likely pathway is validation as Laboratory Developed Test under CLIA, with estimated 2–3 years for regulatory clearance.

Fourth, potential failure modes with real data include encoding errors from non-standard VCF formats or complex variants not handled by preprocessing, reduced accuracy in highly consanguineous populations where long shared haplotypes violate independence assumptions, vulnerability to adversarial variants intentionally designed to cause encoding collisions (requiring attacker knowledge of basis vectors), and batch effects correlating with sensitive attributes that enable inference attacks exploiting technical artifacts rather than biological signal.

Fifth, deployment cost estimates (\$167–\$3,439/month) assume successful validation and regulatory approval. Real-world costs would include personnel (bioinformaticians, system administrators), regulatory compliance (CLIA certification, CAP accreditation), integration with existing systems (EHR interfaces, LIMS connectivity), and ongoing maintenance and updates. These additional costs could substantially exceed infrastructure estimates.

5.4 Ethical Considerations

Strong biometric identification ($D' = 38.43$) could enable surveillance if deployed without appropriate safeguards. Mitigation strategies include rate limiting (1,000 queries/day hard limit), audit logging of all access, ethical use agreements with institutional oversight, and multi-party governance for basis vector management. Open-source implementation prevents vendor lock-in but requires community oversight to prevent misuse. Individual cryptographic control (destroy private key $\rightarrow$ data unrecoverable) aligns with data sovereignty principles but creates challenges for longitudinal research and public health applications requiring persistent identifiers.

5.5 Future Directions

Immediate priorities include validation on real clinical cohorts (1000 Genomes Project: 2,504 individuals; UK Biobank subset), extension to structural variants using long-read sequencing data, validation of mechanistic queries (epistasis detection on quantitative traits: height, BMI, metabolic markers; regulatory network queries on characterized pathways: p53, NF- κ B), and deployment pilot studies with partner hospitals under IRB oversight.

Longer-term research directions include mechanistic applications beyond identification. Three-way epistasis detection currently requires $\sim 10^{18}$ tests for 1M SNPs (computationally intractable); HDC composition through binding operations ($V_1 \odot V_2 \odot V_3$) could reduce complexity to $O(n)$ encoding + $O(k \cdot d)$ per composition, enabling first genome-wide three-way epistasis scan. Protein-DNA binding prediction via compositional queries ($\text{TF}_{HV} \odot \text{Motif}_{HV} \odot \text{Chromatin}_{HV}$) could predict CRISPR off-target effects and design synthetic regulatory circuits. Long-range regulatory inference might detect enhancer-promoter loops from sequence-only data without expensive Hi-C experiments. These mechanistic applications remain theoretical and require extensive validation to determine whether compositional HDC operations capture biological causation.

Integration with blockchain genomics platforms addresses a long-standing technical barrier. Previous platforms (Nebula Genomics, HLTH.network) failed because storing raw genomes on-chain violates privacy regulations while encryption prevents analysis. HDC-encoded genomes (1KB) could be stored on-chain with cryptographic ownership while remaining analyzable via zero-knowledge proofs. This enables direct patient-researcher transactions, immutable provenance tracking, continuous royalty streams via smart contracts, and regulatory compliance (HDC encoding satisfies HIPAA de-identification standard §164.514). Economic model: patients contribute genomes $\rightarrow$ receive ownership tokens $\rightarrow$ earn royalties when data is queried. GenomeVault’s \$0.01–0.10 per query cost enables micro-payments infeasible with traditional blockchain transaction fees.

6 Conclusion

We developed GenomeVault, a privacy-preserving genomic computing platform combining hyperdimensional computing, differential encoding, zero-knowledge proofs, and private information retrieval. Evaluation on a synthetic cohort of 282 subjects demonstrated $\text{AUC} = 1.000$ ($D' = 38.43$) for genetic identification with 5.04ms encoding latency and $178\times$ speedup versus traditional pipelines. Attribute inference attacks achieved only 30% accuracy (baseline: 33%), suggesting effective information concealment under the evaluated threat model. Zero-knowledge proofs generated in 603ms with 100% verification success. Estimated deployment costs range from \$167 to \$3,439 per month depending on scale.

These results on synthetic data indicate that hyperdimensional encoding may enable privacy-preserving genomic analysis with practical performance characteristics. However, validation on real clinical cohorts is essential to confirm these findings and identify failure modes not captured by synthetic data. Extension

to structural variants, regulatory approval, and systematic evaluation of mechanistic applications (epistasis detection, regulatory network queries) represent important next steps. If validated on real data, GenomeVault could enable previously infeasible research applications: global rare disease collaboration with encrypted data sharing, population-scale epistasis studies with reduced computational complexity, real-time clinical genomic matching for emergency diagnostics, and blockchain-based genomic data markets with patient ownership and compensation.

Acknowledgments

[To be added]

Data Availability

Synthetic data generation scripts and validation protocols are available at [. Benchmark results are cryptographically signed and independently verifiable.](#)

Code Availability

Open-source implementation available at [under MIT license.](#)

Competing Interests

The authors declare no competing interests.

Author Contributions

[To be added]

References

- [1] Birney, E. et al. (2021). The future of genomics. *Nature*, 585, 15–19.
- [2] Gymrek, M. et al. (2013). Identifying personal genomes by surname inference. *Science*, 339, 321–324.
- [3] Erlich, Y. & Narayanan, A. (2014). Routes for breaching and protecting genetic privacy. *Nature Reviews Genetics*, 15, 409–421.
- [4] Homer, N. et al. (2008). Resolving individuals contributing trace amounts of DNA to highly complex mixtures. *PLoS Genetics*, 4, e1000167.
- [5] Im, H. K. et al. (2012). On sharing quantitative trait GWAS results in an era of multiple-omics data. *Genetic Epidemiology*, 36, 796–800.
- [6] Kim, M. et al. (2015). Secure genome-wide association analysis using multiparty computation. *Nature Biotechnology*, 33, 1075–1078.
- [7] Chen, F. et al. (2016). HEALER: homomorphic encryption-based secure similarity evaluation. *BMC Medical Genomics*, 9, 46.

- [8] Wang, S. et al. (2017). Genome privacy: challenges, technical approaches to mitigate risk. *Nature Reviews Genetics*, 18, 362–378.
- [9] Kamm, L. et al. (2013). Secure genome-wide association analysis using multiparty computation. *Proceedings on Privacy Enhancing Technologies*, 2013, 91–105.
- [10] Simmons, S. et al. (2016). Enabling privacy-preserving GWASs in heterogeneous populations. *Cell Systems*, 3, 54–61.
- [11] Karunaratne, G. et al. (2021). Hyperdimensional computing: an introduction to computing in distributed representation. *Cognitive Computation*, 13, 1–24.
- [12] Bogatyy, I. & Ishai, Y. (2020). Distributed ORAM for secure multi-server computation. *Theory of Cryptography*, LNCS 12552, 534–565.
- [13] Demmler, D. et al. (2019). Secure evaluation of genomic privacy. *USENIX Security Symposium*, 1093–1110.
- [14] Butler, J. M. (2015). *Advanced Topics in Forensic DNA Typing: Methodology*. Academic Press.
- [15] Phillips, P. J. et al. (2018). Face recognition accuracy of forensic examiners, superrecognizers, and algorithms. *PNAS*, 115, 6171–6176.